

```
import numpy as np
import pandas as pd

#!pip install feature-engine

df=pd.read_csv("/content/housing.csv")

df.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0

```
df.info(memory_usage="deep")

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 2.7 MB

df.describe()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000

scaling is required

Also total rooms and bedrooms are way high for block.....it is given a block in range 600,3000

```
df.isnull().mean()
```

```
longitude      0.000000
latitude       0.000000
housing_median_age  0.000000
total_rooms    0.000000
total_bedrooms  0.010029
population     0.000000
households     0.000000
median_income  0.000000
median_house_value  0.000000
ocean_proximity  0.000000
dtype: float64
```

Since number of null is near 1% and entries are 20640 we can drop null values

```
df["ocean_proximity"].value_counts()
```

```
<1H OCEAN      9136
INLAND         6551
NEAR OCEAN      2658
NEAR BAY        2290
ISLAND           5
Name: ocean_proximity, dtype: int64
```

Double-click (or enter) to edit

```
df.corr()
```

- We can see median_house_value depends highly on median_income which make sense
- Among total_rooms,total_bedrooms,population and households I have choosen *total_rooms as it has higher corr with price* also took latitude

```
housing_median_age    -0.108197    0.011173    1.000000    -0.361262    -0.361262
df=df.dropna()
df=df[df['population']>=600]
df=df[df['population']<=3000]

population    0.000770    0.408785    0.000344    0.057400    0.057400

df.shape

(16271, 10)
```

```
from sklearn.model_selection import train_test_split

y = df['median_house_value']
X = df[['latitude','total_rooms','median_income']]

# For the larger the dataset, the smaller the test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)

X.describe()
```

	latitude	total_rooms	median_income
count	16271.000000	16271.000000	16271.000000
mean	35.569522	2518.339623	3.867861
std	2.113046	1237.362697	1.841316
min	32.540000	121.000000	0.499900
25%	33.930000	1661.500000	2.587500
50%	34.220000	2237.000000	3.550600
75%	37.680000	3087.000000	4.728300
max	41.950000	19107.000000	15.000100

Double-click (or enter) to edit

```
from feature_engine.outliers import Winsorizer
capper = Winsorizer(capping_method='iqr', tail='both')
capper.fit(X_train)
X_train=capper.transform(X_train)
X_test=capper.transform(X_test)

from sklearn.preprocessing import RobustScaler
from feature_engine.wrappers import SklearnTransformerWrapper
scaler = SklearnTransformerWrapper(transformer = RobustScaler())
scaler.fit(X_train)
X_train=scaler.transform(X_train)
```

```
X_test=scaler.transform(X_test)
```

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X_train,y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
pred=model.predict(X_test)
```

Double-click (or enter) to edit

```
from sklearn.metrics import mean_squared_error  
mean_squared_error(y_test,pred)
```

```
6071014206.6652975
```

```
from sklearn.metrics import r2_score  
r2_score(y_test,pred)
```

```
0.5243683694358883
```

